




Exploring **Randomly Wired Neural Networks** for Image Recognition



7th April, 2019
PR12 Paper Review
JinWon Lee
Samsung Electronics

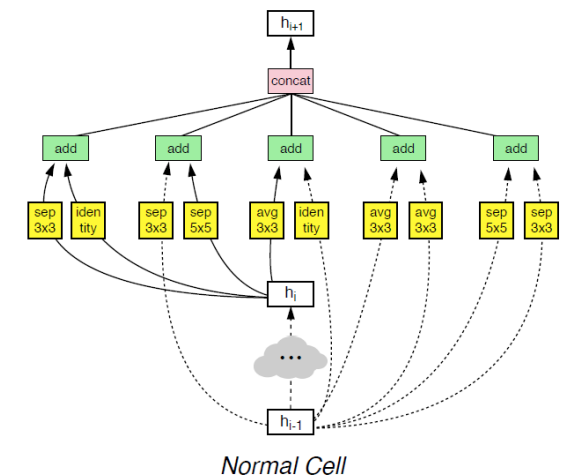

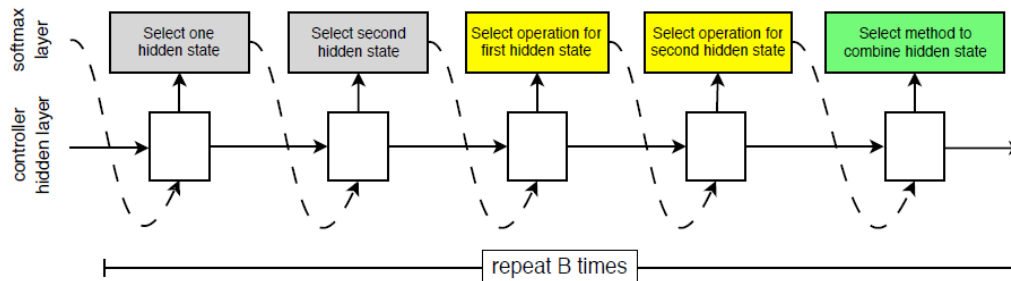
Introduction

- What we call deep learning today descends from **the connectionist approach** to cognitive science.
- **How computational networks are wired is crucial** for building intelligent machines.
- ResNet and DenseNet, that are effective in large part because of **how they are wired**.
- Advancing this trend, **neural architecture search (NAS)** has emerged as a promising direction for **jointly searching wiring patterns and which operations to perform**.



Introduction

- However, like the wiring patterns in ResNet and DenseNet, the NAS network generator is **hand designed** and the space of allowed wiring patterns is **constrained in a small subset of all possible graphs**.
- *What happens if we **loosen this constraint** and **design novel network generators**?*



<NASNet>

Main Topic

- *Design a Network Generator not an Individual Network!*
- The authors suggest a new transition from designing an individual network to designing a network generator may be possible, analogous to how our community have transitioned from designing features to designing a network that learns features.
- Rather than focusing primarily on search with a fixed generator, authors suggest designing new network generators that produce new families of models for searching.

Related Work

- Network Wiring
 - CNN & RNN use **chain-like wiring patterns**.
 - LSTM, Inception, ResNet, and DenseNet **wiring patterns are effective** in general.
- Neural Architecture Search(NAS)
 - Recent research on **NAS mainly focuses on optimization methods**, including RL, progressive, gradient-based, etc.
 - The search space in these NAS works **is largely unchanged** in these works.

Related Work

- Randomly Wired Machines
 - Turing suggested a concept of **unorganized machines**, which is a form of the earliest randomly connected neural networks.
 - Minsky and Rosenblatt also suggested randomly connected machines
- Relation to Neuroscience
 - Turing analogized the unorganized machines to an infant human's cortex.
 - Rosenblatt pointed out that *"at birth, the construction of the most important networks is largely random."*
- Random Graphs in Graph Theory
 - Random graphs are widely studied in graph theory.
 - The definition of the random graph model determines **the prior knowledge encoded in the resulting graphs** and may connect them to naturally occurring phenomena.

Network Generator

- Defining **a network generator** as a mapping g from a parameter space Θ to a space of neural network architecture N .
 - The set N is typically a family of related networks, for example, VGG nets, ResNets, or DenseNets.
- The parameters $\theta \in \Theta$ specify the instantiated network and may contain diverse information.
 - For example, in a ResNet generator, θ can specify the number of stages, number of residual blocks for each stage, depth/width/filter sizes, activation types, etc.

Network Generator


- Stochastic Network Generator
 - The network generator $g(\theta)$ performs *a deterministic mapping*.
 - We can extend g to accept an additional argument s that is the *seed of a pseudo-random number generator* that is used internally by g .
 - We call generators of the form $g(\theta, s)$ *stochastic network generators*.
- NAS from the Network Generator Perspective
 - *The weight matrices of the LSTM are the parameters θ of the generator.*
 - Given the probability distribution conditioned on θ and the seed s , each step samples a construction action(e.g., insert an operator, connect two nodes).
 - NAS is the network generator that is *hand-designed and encodes a prior from human knowledge*.
 - The network space N has been restricted by hand-designed rules. (e.g., 5 nodes in a cell always have input degree 2 and output degree 1)

Randomly Wired Neural Networks

- To investigate how important the generator design is necessary **to study new network generators** that are substantially different from the NAS generator.
- So, authors define network generators that yield networks with random graphs, subject to **different human-specific priors**.
 - They use three classical random graph models – ER, BA, WS models
- Generating General Graphs
 - **Starting by generating a general graph.**
 - **Without restricting** how the graphs correspond to neural networks.
 - Once a graph is obtained, it is mapped to a computable neural networks.
 - The mapping is in itself arbitrary, and authors **intentionally use a simple mapping.**

Randomly Wired Neural Networks

- Edge Operation
 - **The edges are data flow** – a directed edge sends data from one node to another node.
- Node Operation
 - Aggregation
 - The input data to a node are combined via **a weighted sum**; the weights are learnable and positive.
 - Transformation
 - **ReLU-convolution-BN triplet**
 - The same type of convolution is used for all nodes. – 3x3 separable convolution (3x3 depthwise conv \rightarrow 1x1 pointwise conv)
 - Distribution
 - **The same copy of the transformed data** is sent out by the output edges of the node.




Randomly Wired Neural Networks

- Nice Properties of Node Operation
 - Additive aggregation **maintains the same number of output channels as input channels.**
 - Transformed data **can be combined with the data from any other nodes.**
 - Fixing the channel count then **keeps the FLOPs and parameter count unchanged for each node**, regardless of its input and output degrees.
 - Aggregation and distribution are almost parameter-free, regardless of input and output degrees.
 - Also, **the overall FLOPs and parameter count of a graph are roughly proportional to the number of nodes** and nearly independent of the number of edges
 - This enables the comparison of different graphs without inflating/deflating model complexity. **Differences in task performance are therefore reflective of *the properties of the wiring pattern*.**

Randomly Wired Neural Networks

- Input and Output Nodes

- A general graph is not yet a valid neural network.
 - It may have multiple input and output nodes.
- Creating a single extra node that is connected to all original input nodes.
 - Unique input node.
- Similarly, creating a single extra node that is connected to all original output nodes.
 - Unique output node & computing the average(unweighted) from all original output nodes.



Randomly Wired Neural Networks

- Stages

- In image classification in particular, it is common to divide a network into *stages* that progressively down-sample feature maps.
- Simple strategy
 - An entire network consists of multiple stages.
 - One random graph represents one stage, and it is connected to its preceding/succeeding stage by its unique input/output node.
 - For all nodes that are directly connected to the input node, their transformations are modified to have a stride 2.
 - The channel count in a random graph is increased by 2x when going from one stage to the next stage.

<RandWire Architectures>

stage	output	small regime	regular regime
conv ₁	112×112	3×3 conv, $C/2$	
conv ₂	56×56	3×3 conv, C	random wiring $N/2, C$
conv ₃	28×28	random wiring N, C	random wiring $N, 2C$
conv ₄	14×14	random wiring $N, 2C$	random wiring $N, 4C$
conv ₅	7×7	random wiring $N, 4C$	random wiring $N, 8C$
classifier	1×1	1×1 conv, 1280-d global average pool, 1000-d fc, softmax	

Randomly Wired Neural Networks

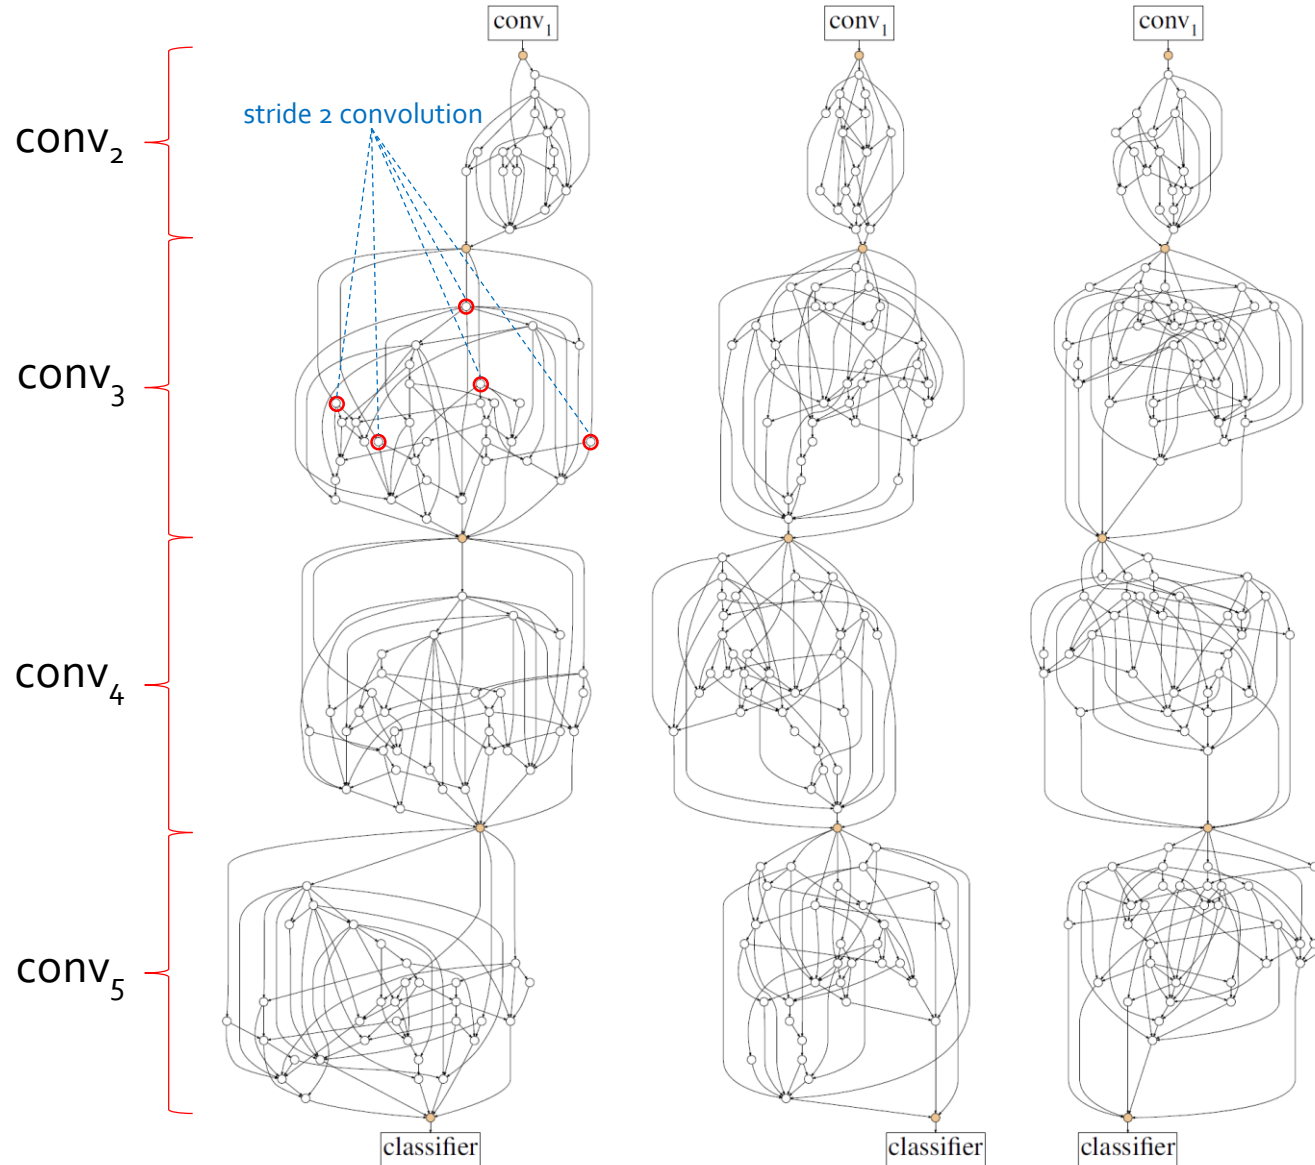





Figure 1. **Randomly wired neural networks** generated by the classical Watts-Strogatz (WS) [50] model: these three instances of random networks achieve (left-to-right) 79.1%, 79.1%, 79.0% classification accuracy on ImageNet under a similar computational budget to ResNet-50, which has 77.1% accuracy.


Random Graph Models

- Erdős-Rényi (ER), 1959.
 - Has N nodes
 - An edge between two nodes is connected with probability P .
 - The ER generation model has only a single parameter P , and is denoted as $ER(P)$.
 - Any graph with N nodes has non-zero probability of being generated by the ER model.




Random Graph Models

- Barabási-Albert (BA), 1999.
 - Generates a random graph by sequentially adding new nodes.
 - The initial state is **M nodes without any edges** ($1 \leq M < N$).
 - **Sequentially adds a new node with M new edges.**
 - For a node to be added, it will be **connected to an existing node v with probability proportional to v's degree.**
 - The new node repeatedly adds **non-duplicate edges** in this way until it has M edges.
 - Any graph generated by BA(M) has exactly $M(N-M)$ edges. **This gives one example on how an underlying prior can be introduced by the graph generator in spite of randomness.**



Random Graph Models

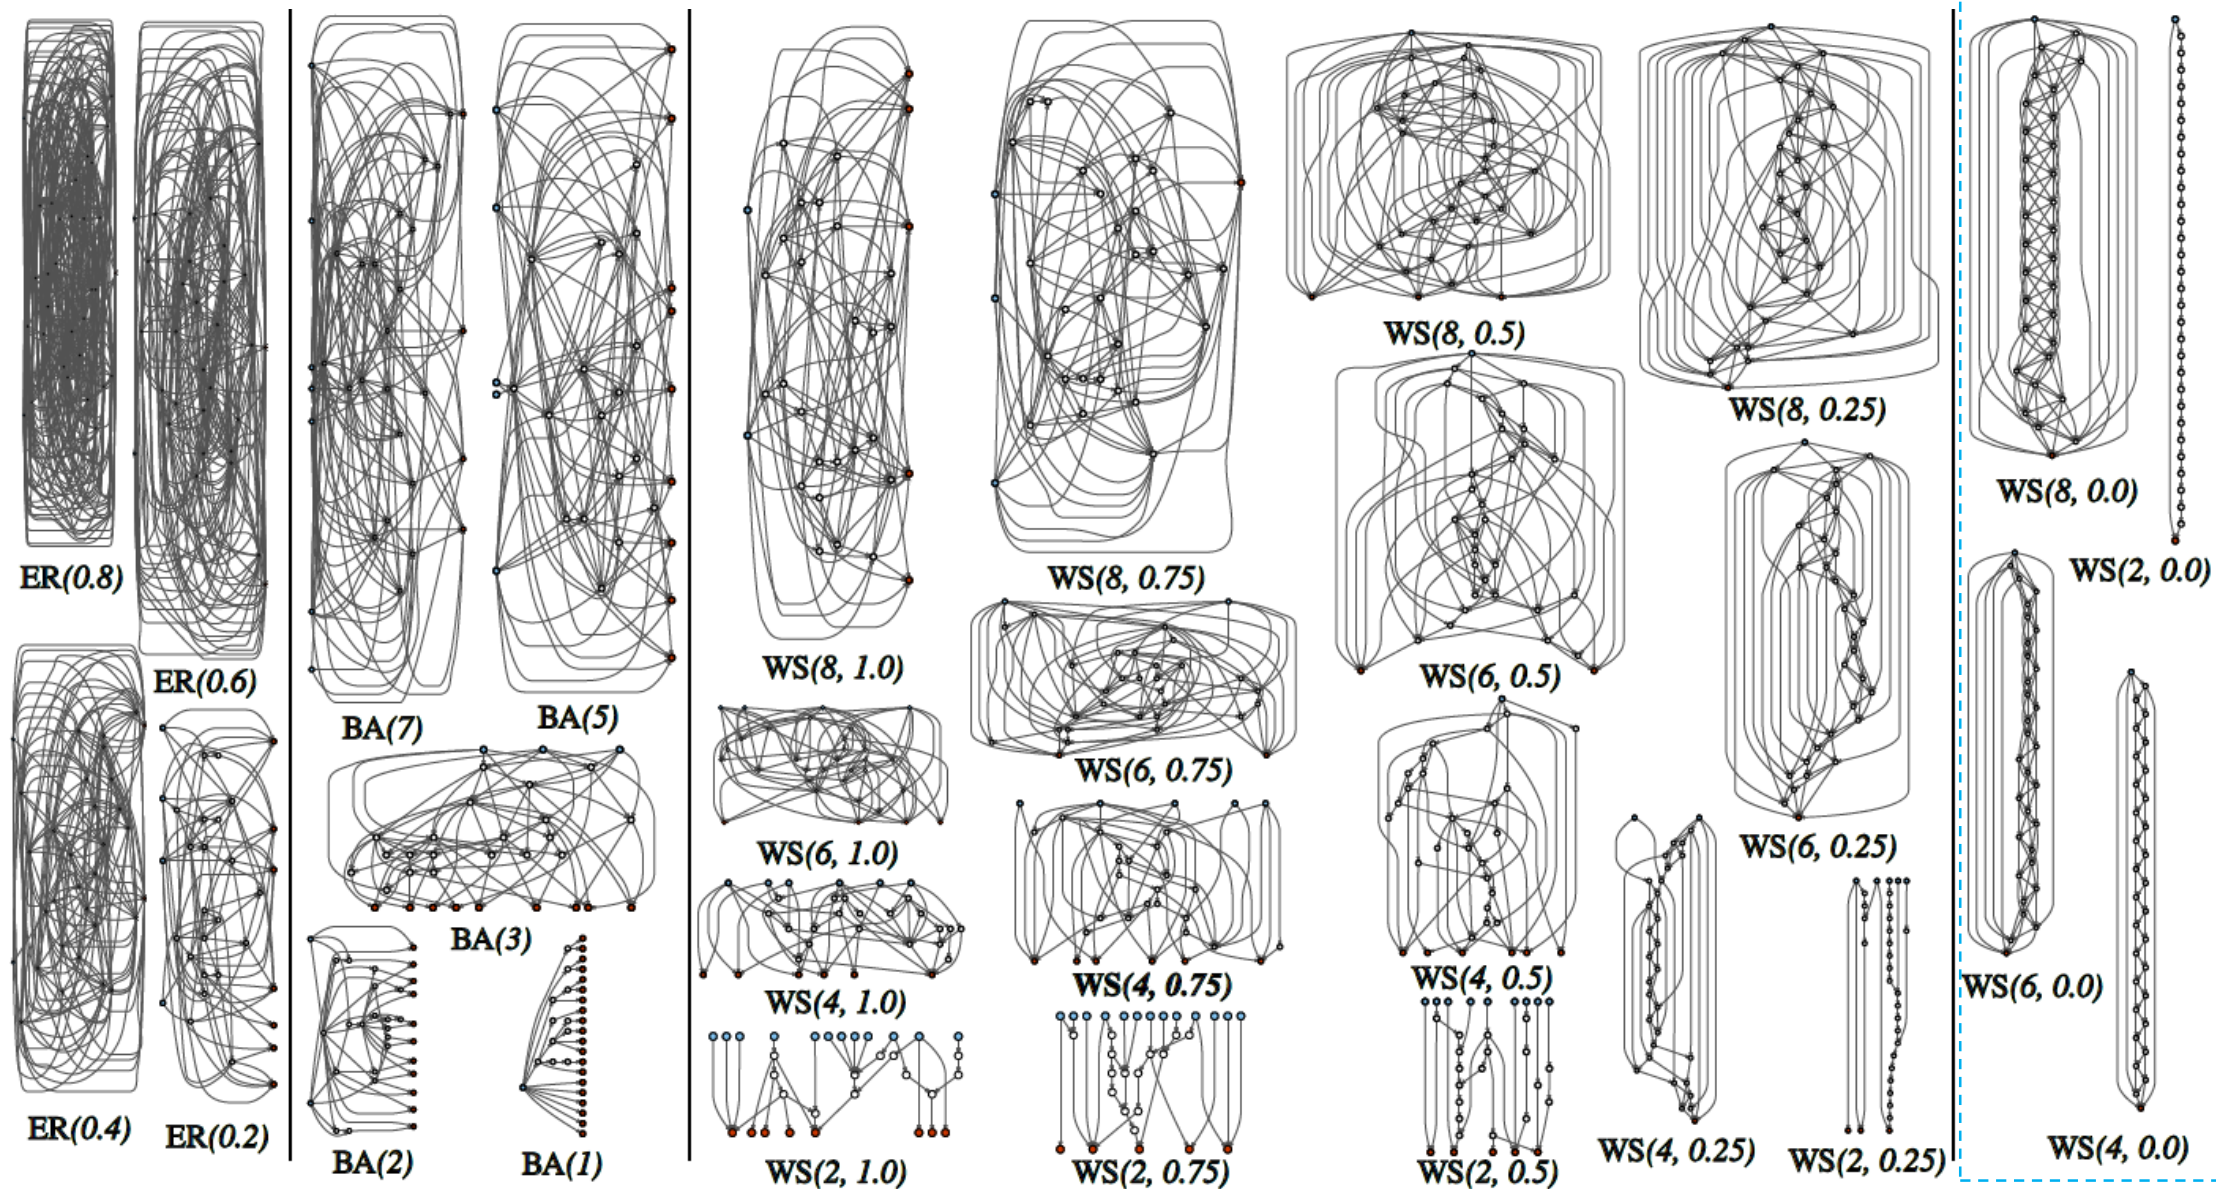
- Watts-Strogatz (WS), 1998.
 - “Small World” model.
 - Stanley Milgram experiment, Erdős number, Bacon number(<http://oracleofbacon.org/>).
 - High clustering, small diameter.
 - N nodes are regularly placed in a ring and each node is connected to its $K=2$ neighbors on both sides – **starts from regular graph**.
 - Then, in a clockwise loop, for every node v , **the edge that connects v to its clockwise i -th next node is rewired with probability P** .
 - “Rewiring” is defined as **uniformly choosing a random node that is not v** and that is not a duplicate edge.



Converting Undirected Graphs into DAGs

- Assign indices to all nodes in a graph and set the direction of every edge as pointing from the smaller-index node to the larger-index one. → no cycle.
- Indexing
 - ER – indices are assigned in a **random order**.
 - BA – **the initial M nodes are assigned indices 1 to M** , and all other nodes are indexed following their **order of adding to the graph**.
 - WS – indices are assigned **sequentially in the clockwise order**.

Visualization of the Random Graphs



Design and Optimization


- Randomly wired neural networks are generated by a **stochastic network generator $g(\theta, s)$** .
- The random graph parameters, $P, M, (K; P)$ in ER, BA, WS respectively, are part of the parameters θ .
- The “optimization” of such a 1- or 2-parameter space is essentially done **by trial-and-error by human designers**. – line/grid search
- The accuracy variation of our networks is small for different seeds s so they perform ***no random search and report mean accuracy*** of multiple random network instances.

Experiments

- Architecture Details
 - A small computation regime – MobileNet & ShuffleNet
 - A regular computation regime – ResNet-50/101
 - N nodes, C channels determine network complexity.
 - N = 32, C = 79 for the small regime.
 - N = 32, C = 109 or 154 for the regular regime.
- Random Seeds
 - Randomly sample 5 network instances, train them from scratch.
 - Report the classification accuracy with “mean \pm std” for all 5 network instances.
- Implementation Details
 - Train for 100 epochs
 - Half-period-cosine shaped learning rate decay and initial learning rate 0.1
 - The weight decay is $5e-5$
 - Momentum 0.9
 - Label smoothing regularization with a coefficient of 0.1

Analysis Experiments

- All random generators provide decent accuracy over all 5 random network instances. – no fails to converge.
- ER, BA, and WS all have certain setting that yield mean accuracy $> 73\%$, with in a $< 1\%$ gap from the best accuracy from WS(4, 0.75)
- The variation among the random network instances is low. (std : 0.2~0.4%)
- Different random generators may have a gap between their mean accuracies. It means that random generator design plays an important role in the accuracy.



Analysis Experiments

- Graph Damage.
 - Graph damage by randomly removing one node or edge.
 - For networks generated by WS, the mean degradation of accuracy is larger when the output degree of the removed node is higher (hub nodes).
 - The accuracy loss is generally decreasing along the x-axis in the right figure (bottom).

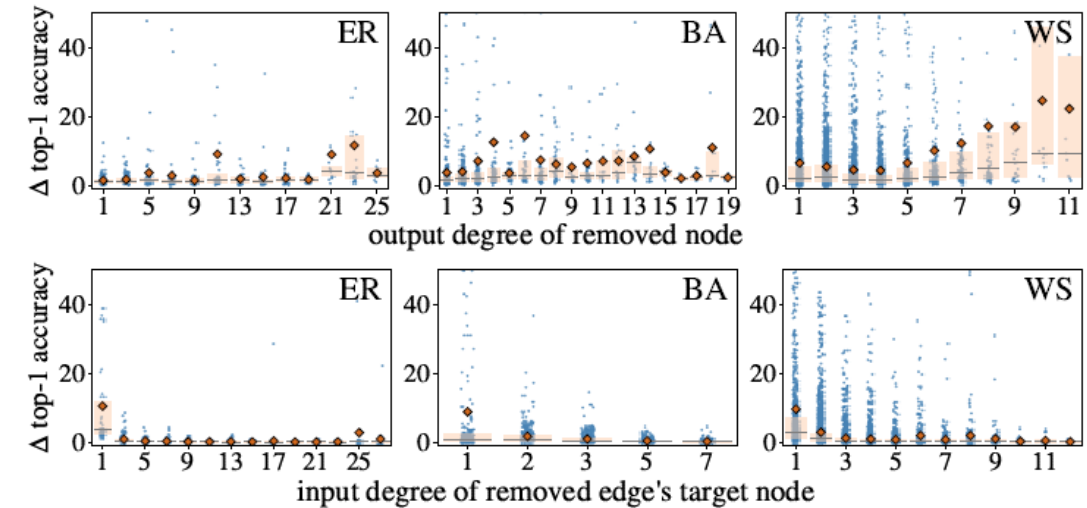
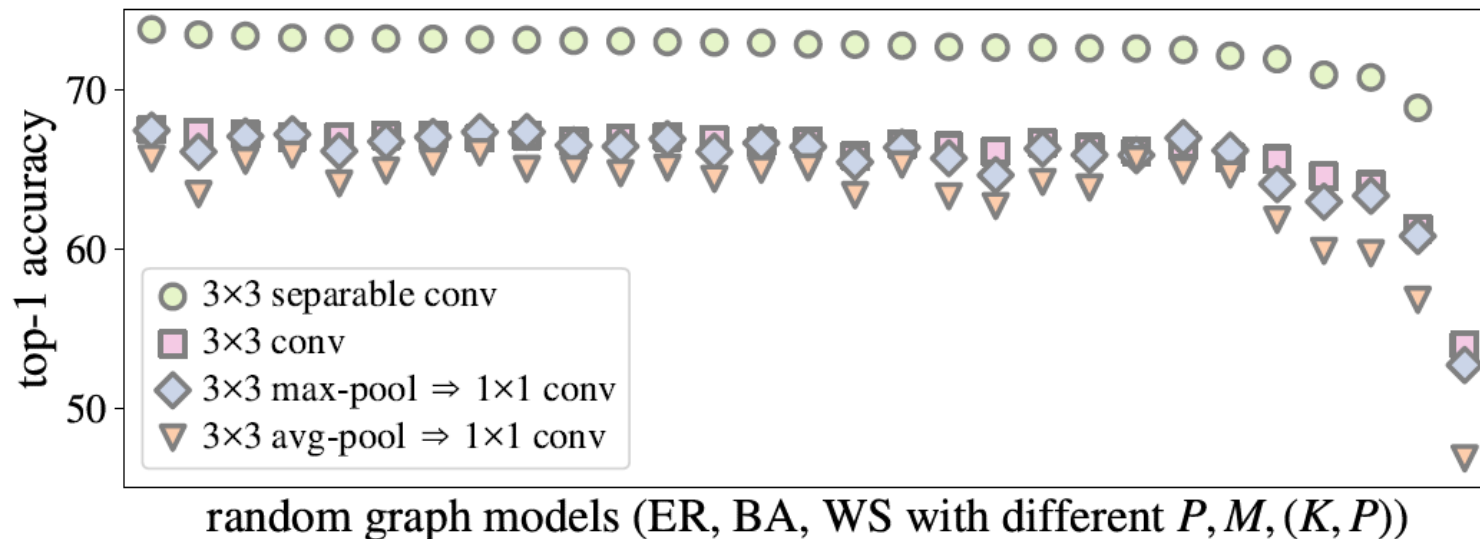


Figure 5. **Graph damage ablation.** We randomly *remove one node* (top) or *remove one edge* (bottom) from a graph after the network is trained, and evaluate the loss (Δ) in accuracy on ImageNet. From left to right are ER, BA, and WS generators. Red circle: *mean*; gray bar: *median*; orange box: *interquartile range*; blue dot: *an individual damaged instance*.

Analysis Experiments

- Node Operations
 - Almost all networks still converge to non-trivial results.
 - The Pearson correlation between any two series in the below figure is 0.91~0.98.
 - This suggests that the network wiring plays a role somewhat orthogonal to the role of the chosen operations.



Comparisons

- Small Computation Regime
 - 250 epochs for fair comparisons.
 - The mean accuracy achieved by RandWire is a competitive result, especially considering that they perform no random search, and use a single operation type for all nodes.

network	top-1 acc.	top-5 acc.	FLOPs (M)	params (M)
MobileNet [15]	70.6	89.5	569	4.2
MobileNet v2 [40]	74.7	-	585	6.9
ShuffleNet [54]	70.9	89.8	524	~5
ShuffleNet v2 [30]	73.7	-	524	~5
NASNet-A [56]	74.0	91.6	564	5.3
NASNet-B [56]	72.8	91.3	488	5.3
NASNet-C [56]	72.5	91.0	558	4.9
Amoeba-A [34]	74.5	92.0	555	5.1
Amoeba-B [34]	74.0	91.5	555	5.3
Amoeba-C [34]	75.7	92.4	570	6.4
PNAS [26]	74.2	91.9	588	5.1
DARTS [27]	73.1	91.0	595	4.9
RandWire-WS	74.7 ± 0.25	92.2 ± 0.15	583 ± 6.2	5.6 ± 0.1

Comparisons

- Regular Computation Regime
 - Use a regularization method inspired by edge removal analysis. Randomly remove one edge whose target node has input degree > 1 with probability of 0.1.
 - Mean accuracies are respectively 1.9% and 1.3% higher than the ResNet-50 and ResNet-101 and are 0.6% higher than the ResNeXt.

network	top-1 acc.	top-5 acc.	FLOPs (B)	params (M)
ResNet-50 [11]	77.1	93.5	4.1	25.6
ResNeXt-50 [52]	78.4	94.0	4.2	25.0
RandWire-WS, $C=109$	79.0 ± 0.17	94.4 ± 0.11	4.0 ± 0.09	31.9 ± 0.66
ResNet-101 [11]	78.8	94.4	7.8	44.6
ResNeXt-101 [52]	79.5	94.6	8.0	44.2
RandWire-WS, $C=154$	80.1 ± 0.19	94.8 ± 0.18	7.9 ± 0.18	61.5 ± 1.32

Comparisons

- Larger Computation

- Same trained networks as regular computation regime, but only increase the test image size to 320x320 without retraining.
- Mean accuracy is 0.7%~1.3% lower than the most accurate NAS results, but use only ~2/3 FLOPs and ~3/4 parameters. Networks are trained for 100 epochs and not on the target image size.

- Object Detection

- The features learned by randomly wired networks can also transfer.

network	test size	epochs	top-1 acc.	top-5 acc.	FLOPs (B)	params (M)
NASNet-A [56]	331 ²	>250	82.7	96.2	23.8	88.9
Amoeba-B [34]	331 ²	>250	82.3	96.1	22.3	84.0
Amoeba-A [34]	331 ²	>250	82.8	96.1	23.1	86.7
PNASNet-5 [26]	331 ²	>250	82.9	96.2	25.0	86.1
RandWire-WS	320 ²	100	81.6 \pm 0.13	95.6 \pm 0.07	16.0 \pm 0.36	61.5 \pm 1.32

backbone	AP	AP ₅₀	AP ₇₅	AP _S	AP _M	AP _L
ResNet-50 [11]	37.1	58.8	39.7	21.9	40.8	47.6
ResNeXt-50 [52]	38.2	60.5	41.3	23.0	41.5	48.8
RandWire-WS, C=109	39.9	61.9	43.3	23.6	43.5	52.7
ResNet-101 [11]	39.8	61.7	43.3	23.7	43.9	51.7
ResNeXt-101 [52]	40.7	62.9	44.5	24.4	44.8	52.7
RandWire-WS, C=154	41.1	63.1	44.6	24.6	45.1	53.0

COCO object detection

Conclusion

- Exploring randomly wired neural networks by three classical random graph models from graph theory.
- The result were surprising: **the mean accuracy of these models is competitive** with hand-designed and optimized from NAS.
- The authors hope that future work **exploring new generator designs may yield new, powerful networks designs.**